

4-sales-prediction-using-python

July 8, 2024

[]: TASK 4 - SALES PREDICTION USING PYTHON

Sales prediction involves forecasting the amount of a product tha customers
↳will purchase,
taking into account various factors such as
"advertising expenditure, target audience segmentation, and advertising
↳platform selection".

In businesses that offer products or services, the role of a Data Scientist is
↳crucial for predicting future sales.

They utilize machine learning techniques in Python to analyze and interpret
↳data,

allowing them to make informed decisions regarding advertising costs.

By leveraging these predictions, businesses can optimize their advertising
↳strategies and maximize sales potential.

Let's embark on the journey of sales prediction using machine learning in
↳Python.

[]:

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
import plotly.io as plio
plio.templates
import plotly.express as px
import plotly.graph_objects as go
from sklearn.model_selection import train_test_split

from sklearn.metrics import mean_squared_error, r2_score

from sklearn.linear_model import LinearRegression, Ridge, Lasso, ElasticNet
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor

from xgboost import XGBRegressor
```

```
import joblib

from warnings import filterwarnings
filterwarnings(action='ignore')
```

```
[2]: data = pd.read_csv(r"C:\Users\divya\OneDrive\Documents\CodSoft_
↳Internship\advertising.csv")
```

```
[3]: data
```

```
[3]:
```

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9
..
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

[200 rows x 4 columns]

```
[4]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   TV           200 non-null    float64
1   Radio        200 non-null    float64
2   Newspaper    200 non-null    float64
3   Sales        200 non-null    float64
dtypes: float64(4)
memory usage: 6.4 KB
```

```
[6]: data.describe()
```

```
[6]:
```

	TV	Radio	Newspaper	Sales
count	200.000000	200.000000	200.000000	200.000000
mean	147.042500	23.264000	30.554000	15.130500
std	85.854236	14.846809	21.778621	5.283892
min	0.700000	0.000000	0.300000	1.600000
25%	74.375000	9.975000	12.750000	11.000000


50%	149.750000	22.900000	25.750000	16.000000
75%	218.825000	36.525000	45.100000	19.050000
max	296.400000	49.600000	114.000000	27.000000

```
[7]: data.duplicated().sum()
```

```
[7]: 0
```

```
[8]: sns.heatmap(data.corr(),annot=True)
```

```
[8]: <Axes: >
```




```
[9]: fig = go.Figure(data=go.Scatter(x=data['TV'], y=data['Sales'], mode='markers',
    ↪marker=dict(color='orange', size=8)))
fig.update_layout(
    title="Scatter Plot of TV(Feature) vs Sales(Target)",
    xaxis_title="TV",
    yaxis_title="Sales"
)
fig.show()
```

```
[10]: fig = go.Figure(data=go.Scatter(x=data['Radio'], y=data['Sales'],
    ↪mode='markers', marker=dict(color='red', size=8)))
fig.update_layout(
    title="Scatter Plot of Radio(Feature) vs Sales(Target)",
    xaxis_title="Radio",
    yaxis_title="Sales"
)
fig.show()
```

```
[11]: fig = go.Figure(data=go.Scatter(x=data['Newspaper'], y=data['Sales'],
    ↪mode='markers', marker=dict(color='blue', size=8)))
fig.update_layout(
    title="Scatter Plot of Newspaper(Feature) vs Sales(Target)",
    xaxis_title="Newspaper",
    yaxis_title="Sales"
)
fig.show()
```

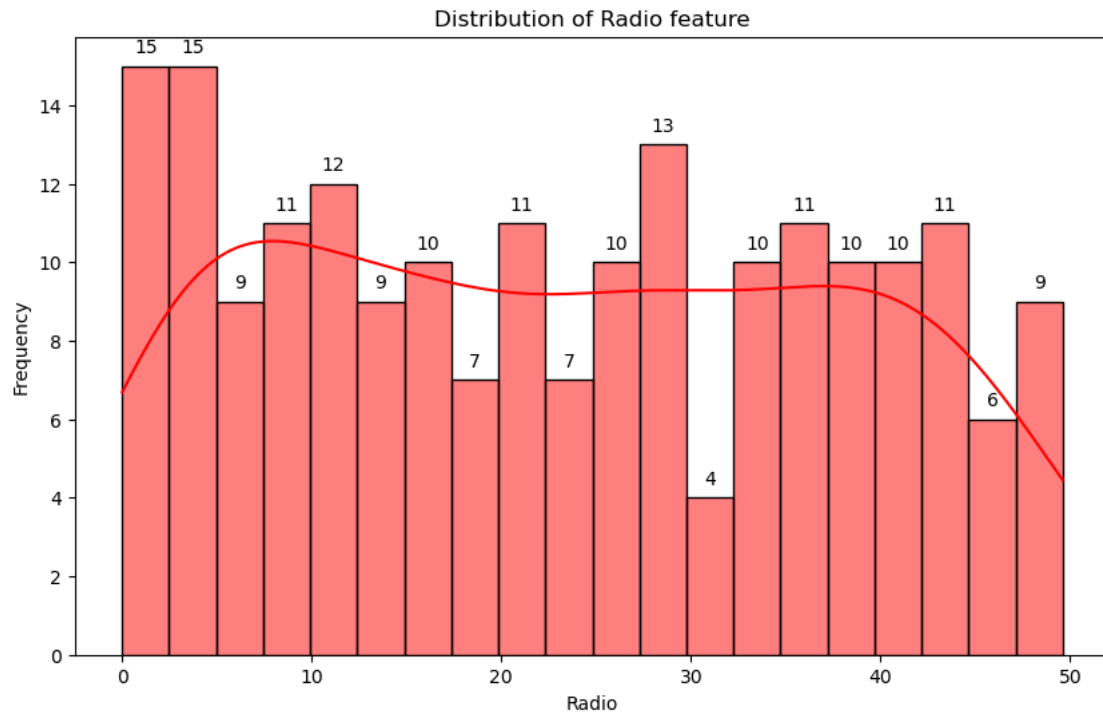
[]: 1. DISTRIBUTION OF TV FEATURE

```
[12]: plt.figure(figsize=(10, 6))
ax = sns.histplot(data['TV'], bins=20, kde=True, color='orange')
plt.xlabel('TV')
plt.ylabel('Frequency')
plt.title('Distribution of TV feature')
for p in ax.patches:
    ax.annotate(f'{p.get_height()}', (p.get_x() + p.get_width() / 2., p.
    ↪get_height()),
                ha='center', va='center', fontsize=10, color='black',
    ↪xytext=(0, 10),
                textcoords='offset points')
plt.show()
```



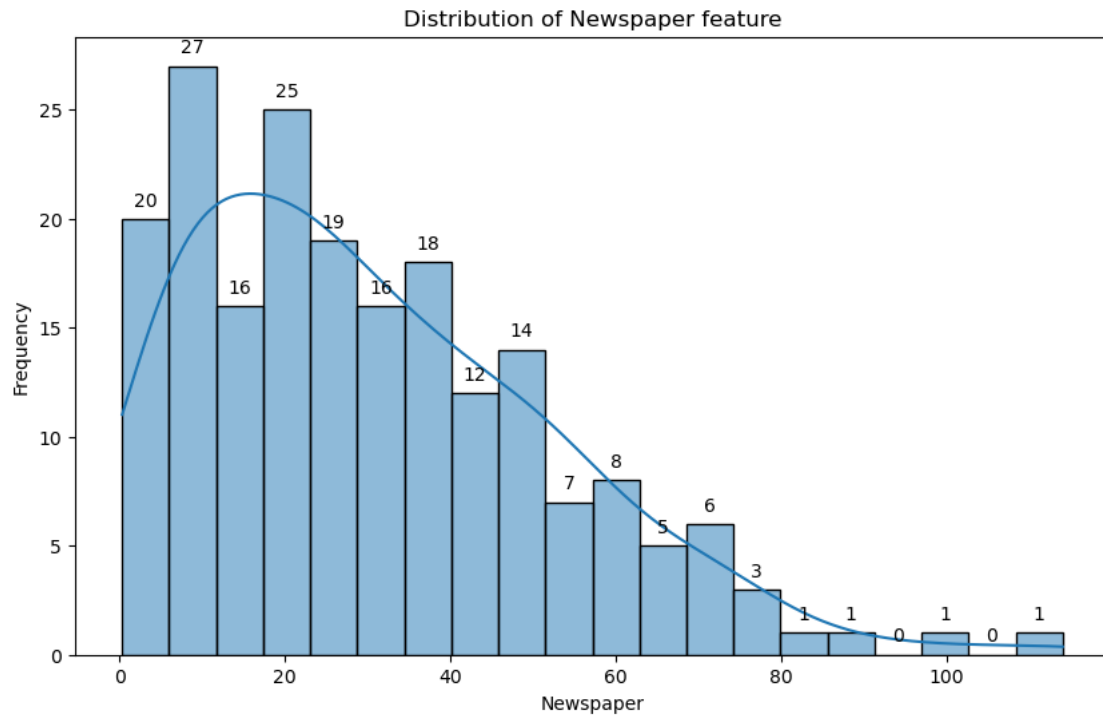
[]: 2. DISTRIBUTION OF RADIO FEATURE

```
[13]: plt.figure(figsize=(10, 6))
ax = sns.histplot(data['Radio'], bins=20, kde=True, color='red')
plt.xlabel('Radio')
plt.ylabel('Frequency')
plt.title('Distribution of Radio feature')
for p in ax.patches:
    ax.annotate(f'{p.get_height()}', (p.get_x() + p.get_width() / 2., p.
    ↪get_height()),
                ha='center', va='center', fontsize=10, color='black',
    ↪xytext=(0, 10),
                textcoords='offset points')
plt.show()
```



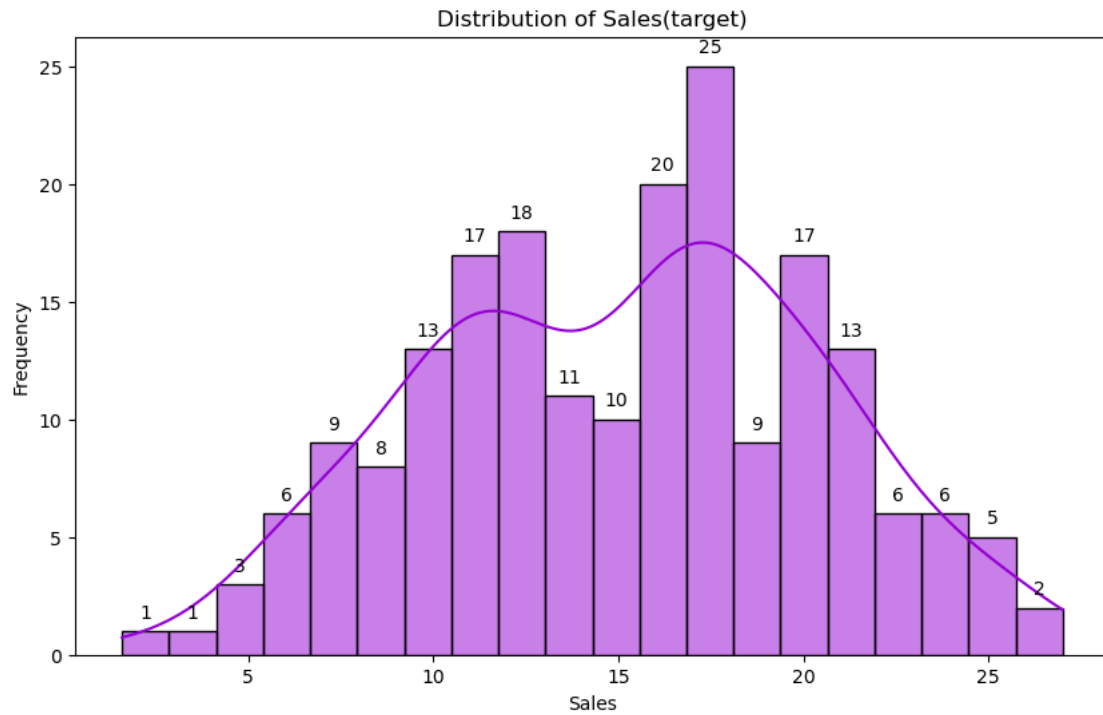
[]: 3. DISTRIBUTION OF NEWSPAPER FEATURE

```
[14]: plt.figure(figsize=(10, 6))
ax = sns.histplot(data['Newspaper'], bins=20, kde=True)
plt.xlabel('Newspaper')
plt.ylabel('Frequency')
plt.title('Distribution of Newspaper feature')
for p in ax.patches:
    ax.annotate(f'{p.get_height()}', (p.get_x() + p.get_width() / 2., p.
    ↪get_height()),
                ha='center', va='center', fontsize=10, color='black',
    ↪xytext=(0, 10),
                textcoords='offset points')
plt.show()
```



[]: 4. DISTRIBUTION OF SALES

```
[15]: plt.figure(figsize=(10, 6))
ax = sns.histplot(data['Sales'], bins=20, kde=True, color='darkviolet')
plt.xlabel('Sales')
plt.ylabel('Frequency')
plt.title('Distribution of Sales(target)')
for p in ax.patches:
    ax.annotate(f'{p.get_height()}', (p.get_x() + p.get_width() / 2., p.
    ↪get_height()),
                ha='center', va='center', fontsize=10, color='black',
    ↪xytext=(0, 10),
                textcoords='offset points')
```



```
[16]: x = data.iloc[:, :3]
      y = data.iloc[:, 3:]
```

```
[17]: x
```

```
[17]:
```

	TV	Radio	Newspaper
0	230.1	37.8	69.2
1	44.5	39.3	45.1
2	17.2	45.9	69.3
3	151.5	41.3	58.5
4	180.8	10.8	58.4
..
195	38.2	3.7	13.8
196	94.2	4.9	8.1
197	177.0	9.3	6.4
198	283.6	42.0	66.2
199	232.1	8.6	8.7

[200 rows x 3 columns]

```
[18]: y
```

```
[18]:
```

	Sales
0	22.1


```
1      10.4
2      12.0
3      16.5
4      17.9
...
195     7.6
196    14.0
197    14.8
198    25.5
199    18.4
```

```
[200 rows x 1 columns]
```

```
[19]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3)
```

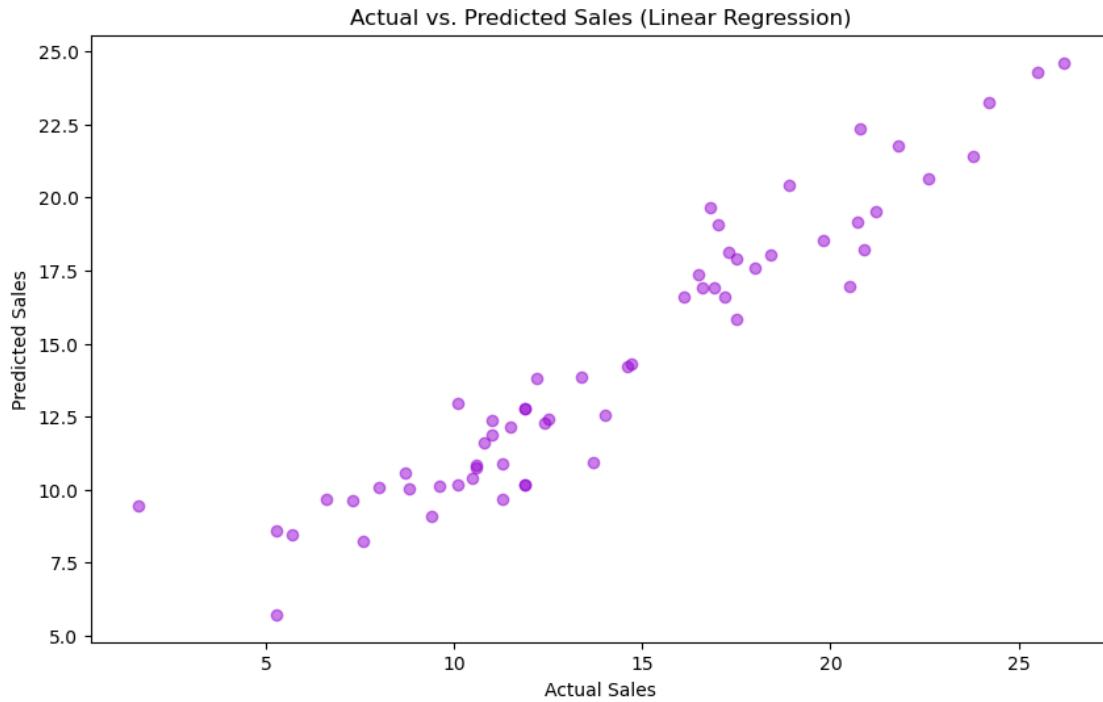
```
[20]: model_1 = LinearRegression()
model_1.fit(x_train, y_train)
y_pred_1 = model_1.predict(x_test)
```

```
[21]: mse = mean_squared_error(y_test, y_pred_1)
r2 = r2_score(y_test, y_pred_1)
print("Mean Square Error is :", mse)
print("R-Squared score is :", r2)
```

```
Mean Square Error is : 3.4038325522795776
```

```
R-Squared score is : 0.8875641116008756
```

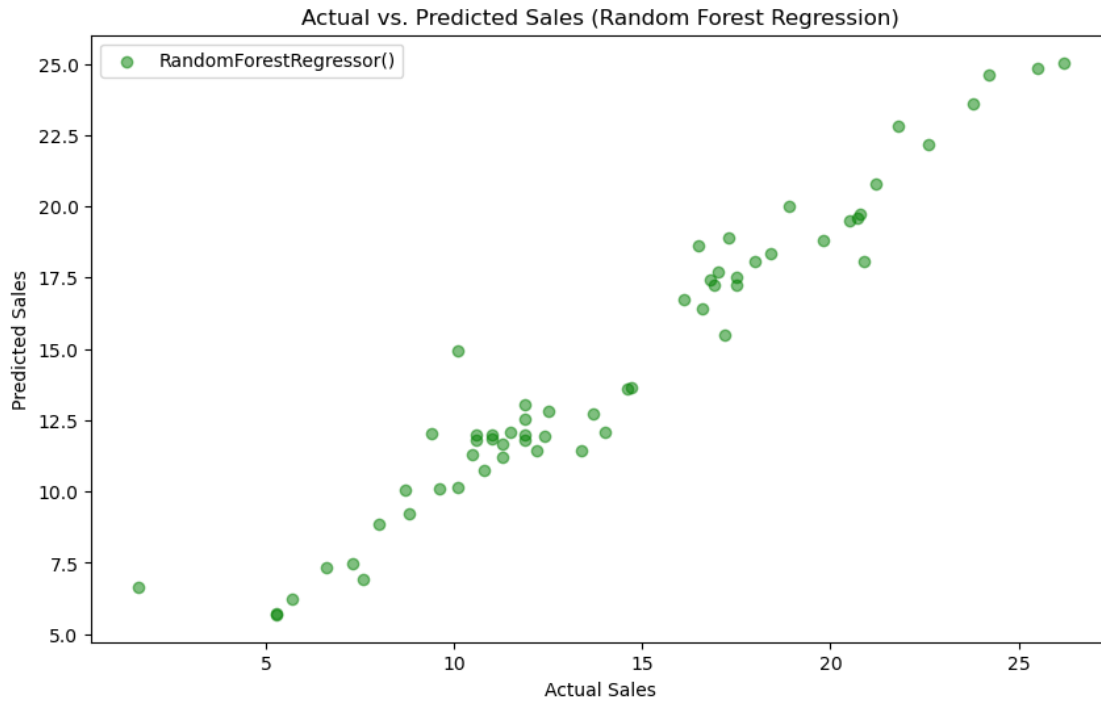
```
[22]: plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred_1, alpha=0.5, color='darkviolet')
plt.xlabel('Actual Sales')
plt.ylabel('Predicted Sales')
plt.title('Actual vs. Predicted Sales (Linear Regression)')
plt.show()
```



```
[23]: model_2 = RandomForestRegressor()
model_2.fit(x_train, y_train)
y_pred_2 = model_2.predict(x_test)
mse = mean_squared_error(y_test, y_pred_2)
r2 = r2_score(y_test, y_pred_2)
print("Mean Square Error is :", mse)
print("R-Squared score is :", r2)
```

```
Mean Square Error is : 1.8022122666666631
R-Squared score is : 0.940469064158059
```

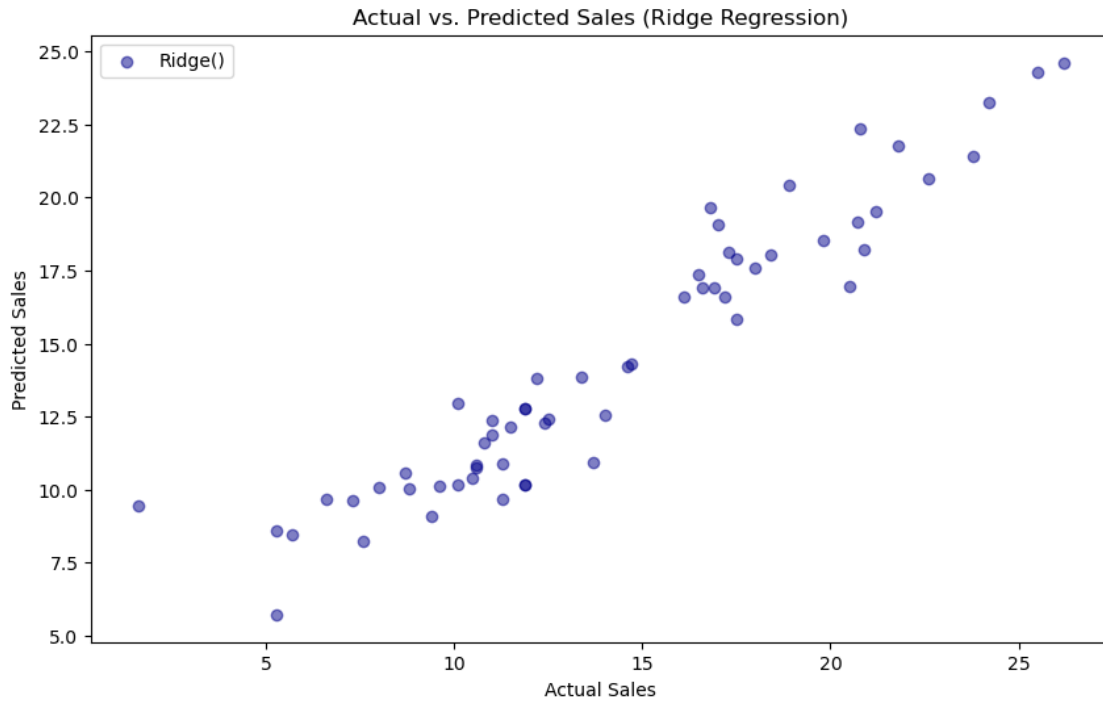
```
[24]: plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred_2, label=model_2, alpha=0.5, color='green')
plt.xlabel('Actual Sales')
plt.ylabel('Predicted Sales')
plt.title('Actual vs. Predicted Sales (Random Forest Regression)')
plt.legend()
plt.show()
```



```
[25]: model_3 = Ridge()
model_3.fit(x_train, y_train)
y_pred_3 = model_3.predict(x_test)
mse = mean_squared_error(y_test, y_pred_3)
r2 = r2_score(y_test, y_pred_3)
print("Mean Square Error is :", mse)
print("R-Squared score is :", r2)
```

Mean Square Error is : 3.4038169687278437
R-Squared score is : 0.8875646263590068

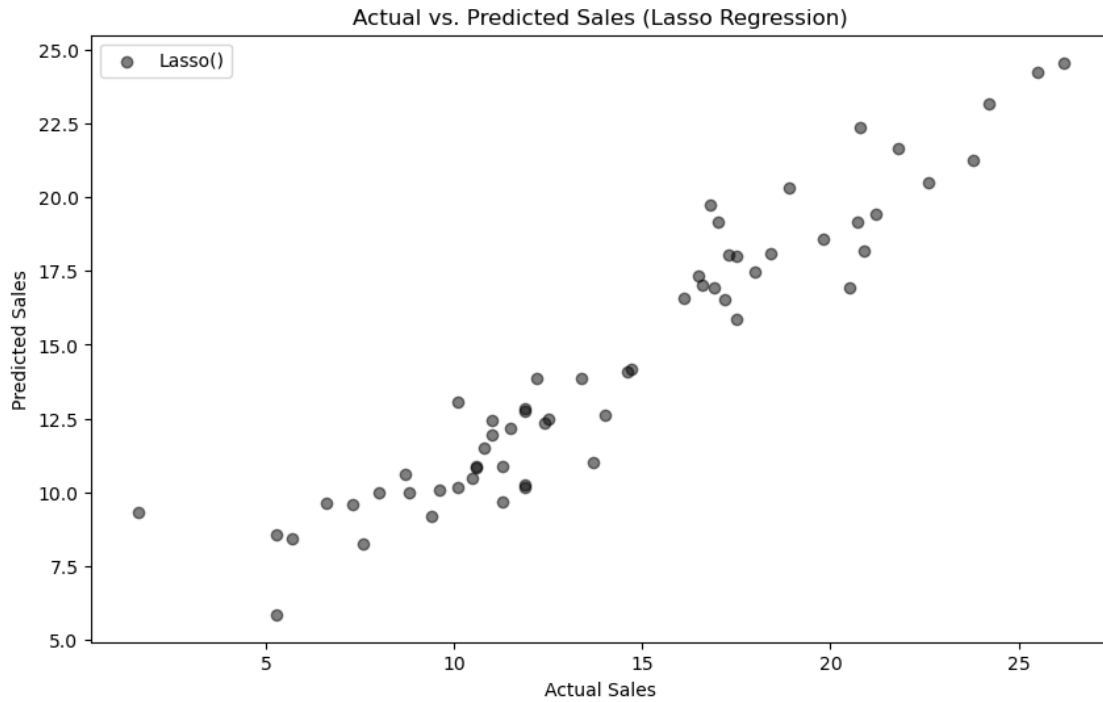
```
[26]: plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred_3, label=model_3, alpha=0.5, color='darkblue')
plt.xlabel('Actual Sales')
plt.ylabel('Predicted Sales')
plt.title('Actual vs. Predicted Sales (Ridge Regression)')
plt.legend()
plt.show()
```



```
[27]: model_4 = Lasso()
model_4.fit(x_train, y_train)
y_pred_4 = model_4.predict(x_test)
mse = mean_squared_error(y_test, y_pred_4)
r2 = r2_score(y_test, y_pred_4)
print("Mean Square Error is :", mse)
print("R-Squared score is :", r2)
```

Mean Square Error is : 3.3925415015877234
R-Squared score is : 0.887937079217819

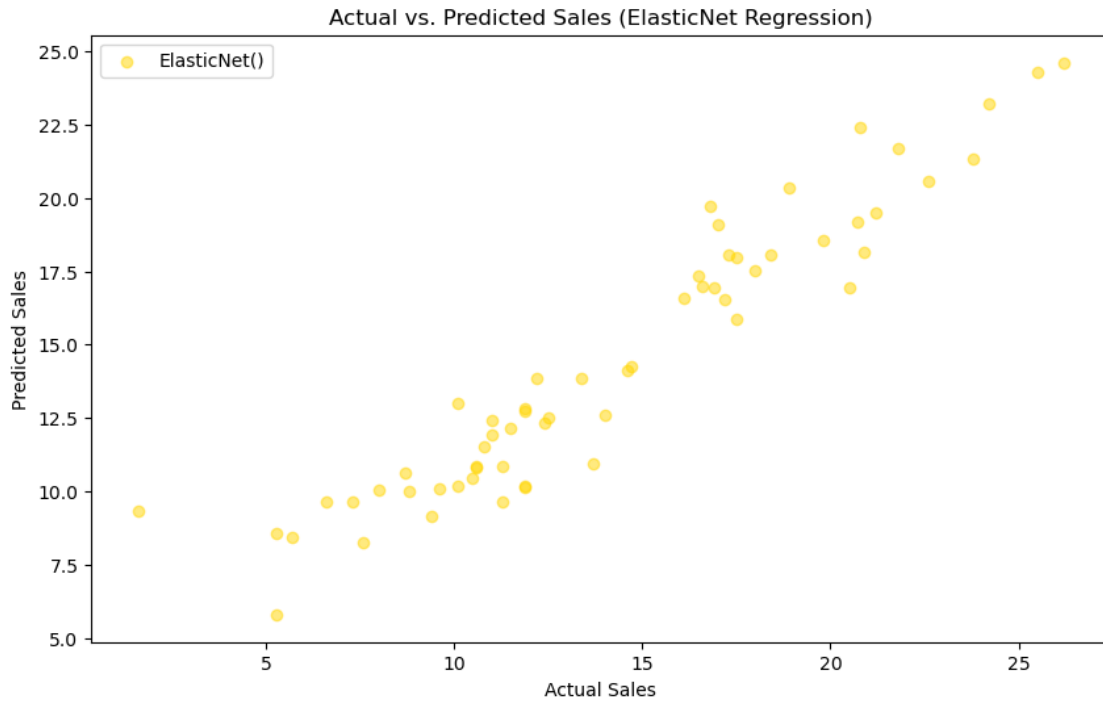
```
[28]: plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred_4, label=model_4, alpha=0.5, color='black')
plt.xlabel('Actual Sales')
plt.ylabel('Predicted Sales')
plt.title('Actual vs. Predicted Sales (Lasso Regression)')
plt.legend()
plt.show()
```



```
[29]: model_5 = ElasticNet()
model_5.fit(x_train, y_train)
y_pred_5 = model_5.predict(x_test)
mse = mean_squared_error(y_test, y_pred_5)
r2 = r2_score(y_test, y_pred_5)
print("Mean Square Error is :", mse)
print("R-Squared score is :", r2)
```

Mean Square Error is : 3.390505094047334
R-Squared score is : 0.8880043461257616

```
[30]: plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred_5, label=model_5, alpha=0.5, color='gold')
plt.xlabel('Actual Sales')
plt.ylabel('Predicted Sales')
plt.title('Actual vs. Predicted Sales (ElasticNet Regression)')
plt.legend()
plt.show()
```



```
[31]: model_6 = GradientBoostingRegressor()
model_6.fit(x_train, y_train)
y_pred_6 = model_6.predict(x_test)
mse = mean_squared_error(y_test, y_pred_6)
r2 = r2_score(y_test, y_pred_6)
print("Mean Square Error is :", mse)
print("R-Squared score is :", r2)
```

Mean Square Error is : 2.0098563088322723
R-Squared score is : 0.9336101361722977


```
[32]: plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred_6, label=model_6, alpha=0.5, color='red')
plt.xlabel('Actual Sales')
plt.ylabel('Predicted Sales')
plt.title('Actual vs. Predicted Sales (Gradient Boosting Regression)')
plt.legend()
plt.show()
```



```
[33]: model_7 = XGBRegressor()
model_7.fit(x_train, y_train)
y_pred_7 = model_7.predict(x_test)
mse = mean_squared_error(y_test, y_pred_7)
r2 = r2_score(y_test, y_pred_7)
print("Mean Square Error is :", mse)
print("R-Squared score is :", r2)
```

Mean Square Error is : 1.82916503673819
R-Squared score is : 0.9395787563649274

```
[34]: plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred_7, label='XGBoost Regression Model', alpha=0.5,
           color='brown')
plt.xlabel('Actual Sales')
plt.ylabel('Predicted Sales')
plt.title('Actual vs. Predicted Sales (XGBoost Regression)')
plt.legend()
plt.show()
```



```
[35]: model_r2_scores = {
    "Linear Regression Model": r2_score(y_test, y_pred_1),

    "Random Forest Regression Model": r2_score(y_test, y_pred_2),

    "Ridge Regression Model": r2_score(y_test, y_pred_3),

    "Lasso Regression Model": r2_score(y_test, y_pred_4),

    "ElasticNet Regression Model": r2_score(y_test, y_pred_5),

    "Gradient Boosting Regression Model": r2_score(y_test, y_pred_6),

    "XGBoost Regression Model": r2_score(y_test, y_pred_7)
}
model_r2_scores
```

```
[35]: {'Linear Regression Model': 0.8875641116008756,
'Random Forest Regression Model': 0.940469064158059,
'Ridge Regression Model': 0.8875646263590068,
'Lasso Regression Model': 0.887937079217819,
'ElasticNet Regression Model': 0.8880043461257616,
'Gradient Boosting Regression Model': 0.9336101361722977,
'XGBoost Regression Model': 0.9395787563649274}
```



```
[36]: best_model_name = max(model_r2_scores, key=model_r2_scores.get)
      best_r2_score = model_r2_scores[best_model_name]

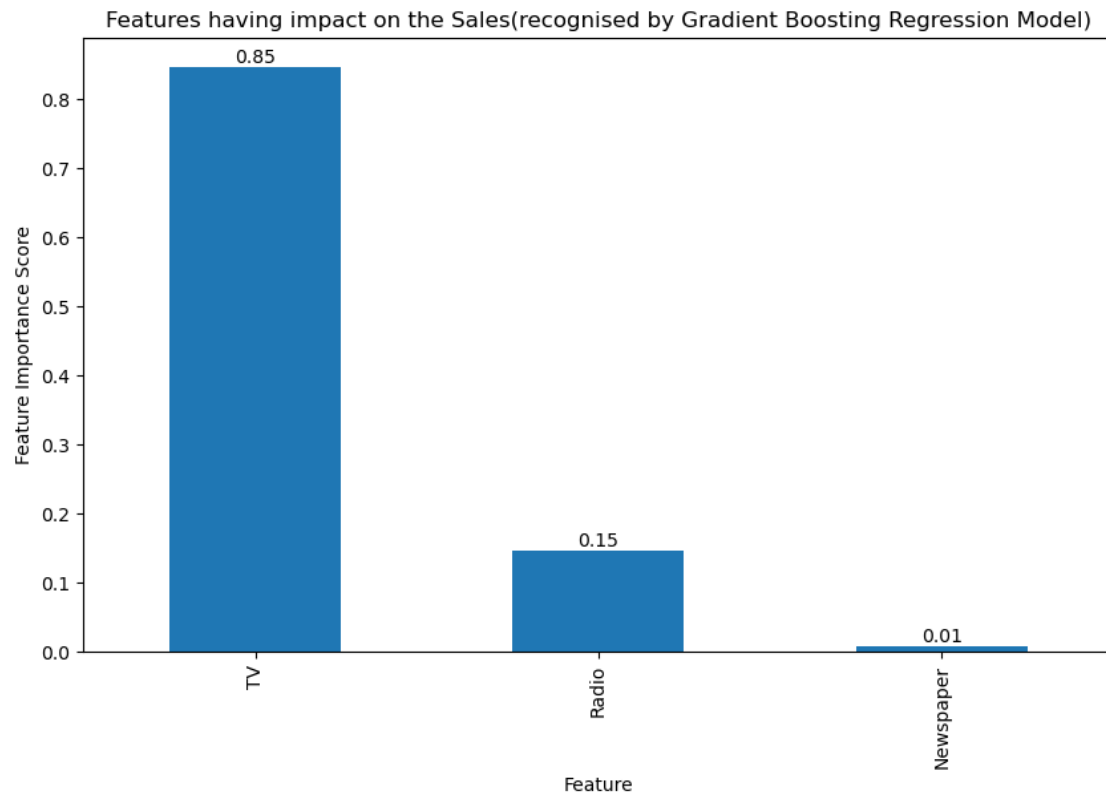
      print(f"Best Performing Model is {best_model_name} with an R^2 score of_
            ↪{best_r2_score}")
```

Best Performing Model is Random Forest Regression Model with an R^2 score of 0.940469064158059

```
[37]: final_model = model_6
      joblib.dump(final_model, 'gradient_boosting_model.pkl')
```

```
[37]: ['gradient_boosting_model.pkl']
```

```
[38]: feature_importances = pd.Series(final_model.feature_importances_, index=x.
            ↪columns)
      plt.figure(figsize=(10, 6))
      features = feature_importances
      features.plot(kind='bar')
      plt.xlabel('Feature')
      plt.ylabel('Feature Importance Score')
      plt.title('Features having impact on the Sales(recognised by Gradient Boosting_
            ↪Regression Model)')
      for index, value in enumerate(features):
          plt.text(index, value, f'{value:.2f}', ha='center', va='bottom')
      plt.show()
```



[]: